



# CREDIT FRAUD DETECTION

SC1015 PRESENTATION

**By:** Jerald (U2221612D)  
Bryan (U2222849J)  
Jonathan (U2223191G)

# TABLE OF CONTENTS

**01**

## DATA SET USED

Motivations

**02**

## DATA RELATED

Visualisation and Cleaning

**03**

## MACHINE LEARNING

Decision Tree, Random Forest,  
K-Nearest Neighbours, Naive Bayes

**04**

## OUTCOME

Our Thoughts



**DATA SET**

# MOTIVATION

- Cashless payments are **increasingly common** with the push of such technology
- To **minimise** the **risk** of lending money to customers who may not pay back their loans.
- To understand which **factors** affect people from defaulting credit card loans.




PAYNOW




# 73%

of Singaporeans own at least 1 credit card in 2023

# DATA SET



application\_data  
columns\_description  
previous\_application



## FEATURE SELECTION

**19**/124 Columns in application\_data

TARGET | NAME\_CONTRACT\_TYPE |  
FLAG\_OWN\_CAR | FLAG\_OWN\_REALTY |  
AMT\_CREDIT | AMT\_GOODS\_PRICE, etc.

## PROBLEM DEFINITION


"To find the **highest correlation** of credit card defaulters and to **predict** loan default risk of defaulters."



# **DATA CLEANING**



# SPLITTING THE DATA



## NUMERICAL

- Finding **NULL** values within dataset
- Fill in missing values with **MEDIAN**

```
[ ] refined_app.isnull().sum()
```

```
SK_ID_CURR      0
TARGET          0
NAME_CONTRACT_TYPE 0
FLAG_OWN_CAR    0
FLAG_OWN_REALTY 0
CNT_CHILDREN    0
AMT_INCOME_TOTAL 0
AMT_CREDIT      0
```

```
AMT_ANNUITY      12
AMT_GOODS_PRICE  278
NAME_TYPE_SUITE  1292
NAME_INCOME_TYPE 0
DAYS_BIRTH       0
DAYS_EMPLOYED    0
OCCUPATION_TYPE  96391
OBS_30_CNT_SOCIAL_CIRCLE 1021
DEF_30_CNT_SOCIAL_CIRCLE 1021
OBS_60_CNT_SOCIAL_CIRCLE 1021
DEF_60_CNT_SOCIAL_CIRCLE 1021
dtype: int64
```

## CATEGORICAL



- Finding **empty values** within dataset
- Fill in missing values with **NULL**


```
.isnull().sum()
```

# FILLING MISSING VALUES

## NUMERICAL

6/13 columns with missing data

```
def median_impute(df,col):  
    return df[col].fillna(df[col].median())  
  
refined_num['AMT_ANNUITY'] = median_impute(refined_num,'AMT_ANNUITY')
```



Filling in with **MEDIAN**

## CATEGORICAL

2/6 columns with missing data

**NAME\_TYPE\_SUITE**: Fill with “*Unaccompanied*”

**OCCUPATION\_TYPE**: Fill with “*NILL*”

**.fillna**

# REMOVING DUPLICATES

- Ensure each record represents a unique observation
- Affects data analysis:
  - Skewness of results of statistical analysis
  - Increasing size of data set unnecessarily
  - Overfitting in machine learning



```
[ ] df = refined_app.copy()
    df.drop_duplicates(inplace=True)
    print('The amount of frauds in df before dropping duplicates:', len(refined_app[refined_app['TARGET'] == 1]))
    print('The amount of frauds in df after dropping duplicates:', len(df[df['TARGET'] == 1]))
```

```
The amount of frauds in df before dropping duplicates: 24825
The amount of frauds in df after dropping duplicates: 24825
```



# DATA VISUALISATION

# NUMERICAL DATA

**Target** (whether individuals defaulted)

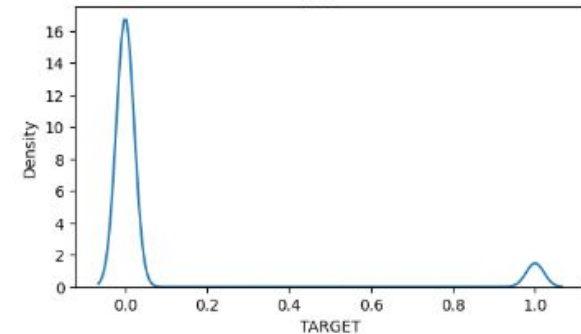


Heavily skewed - more 0s than 1s

## Repercussions:

1. Biased model
2. Poor performance of minority class
3. Misleading evaluation metrics

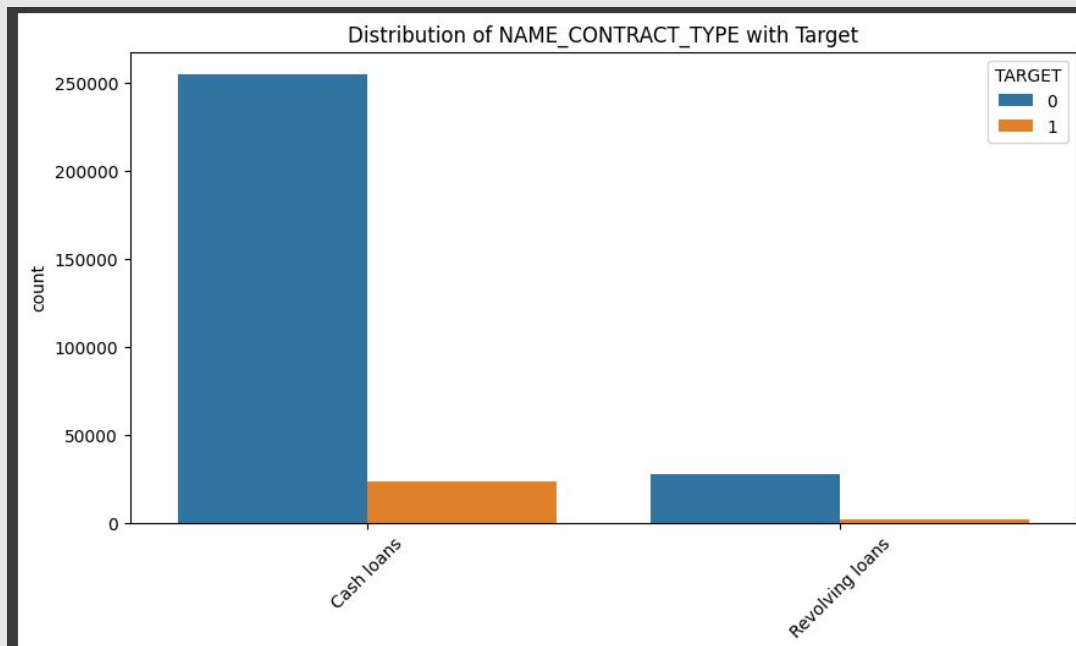
## Possible solutions:

1. Undersampling of majority class
2. Oversampling of minority class (SMOTE)
3. Cost-sensitive learning
4. Ensemble method



# CATEGORICAL DATA

## UNDERSAMPLING



# CORRELATION (RAW DATA)



**TARGET**

**DAYS EMPLOYED**  
**DAYS BIRTH**

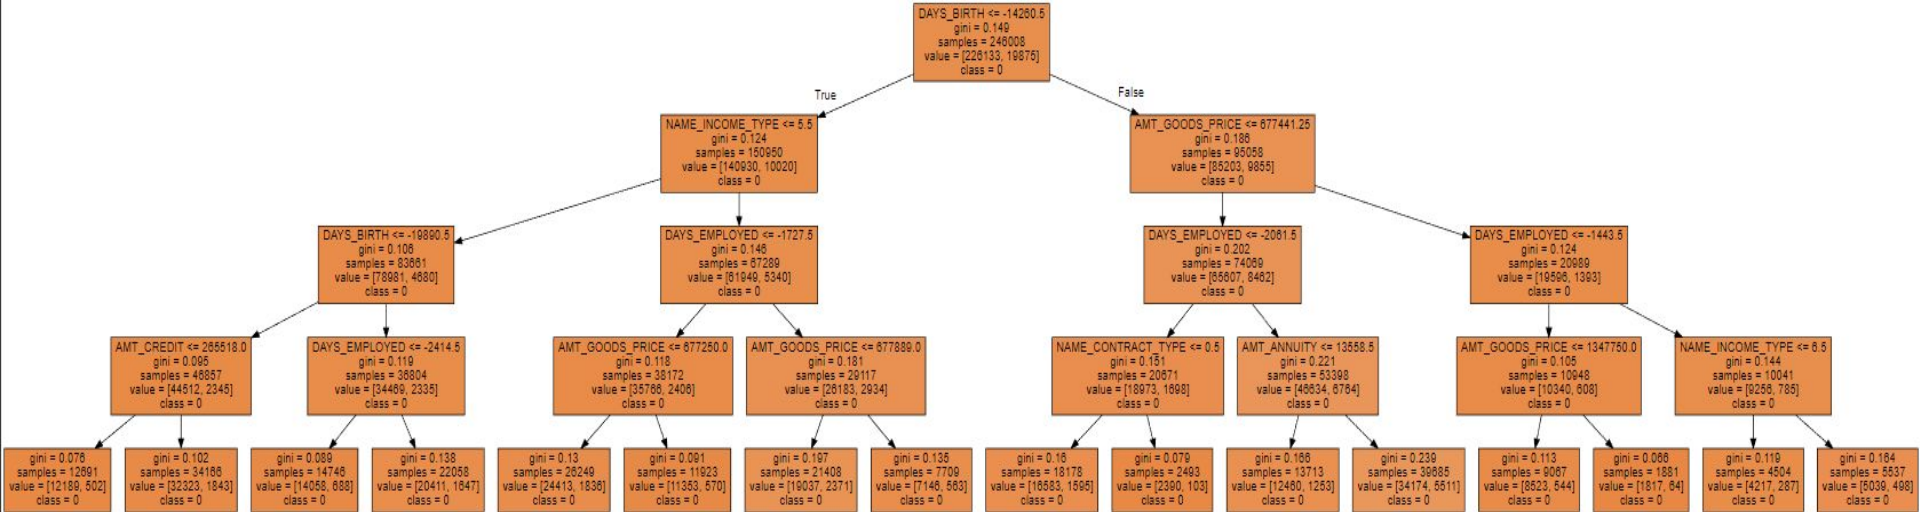
```
DAYS_EMPLOYED      0.071695
DAYS_BIRTH          0.063464
AMT_GOODS_PRICE    -0.029617
AMT_INCOME_TOTAL   -0.024722
AMT_CREDIT          -0.020198
AMT_ANNUITY         0.015305
CNT_CHILDREN        0.010850
OBS_60_CNT_SOCIAL_CIRCLE -0.009135
SK_ID_CURR          0.007665
OBS_30_CNT_SOCIAL_CIRCLE -0.007648
DEF_30_CNT_SOCIAL_CIRCLE      NaN
DEF_60_CNT_SOCIAL_CIRCLE      NaN
Name: TARGET, dtype: float64
The most correlated factor with TARGET is 'DAYS_EMPLOYED', with a correlation value of 0.0717.
Number of frauds (TARGET = 1) after dropping outliers: 1597
```

A black icon on the left side of the text, depicting a stylized brain with circuit-like lines extending from its right side, symbolizing machine learning or artificial intelligence.

# **MACHINE LEARNING**



# DECISION TREE



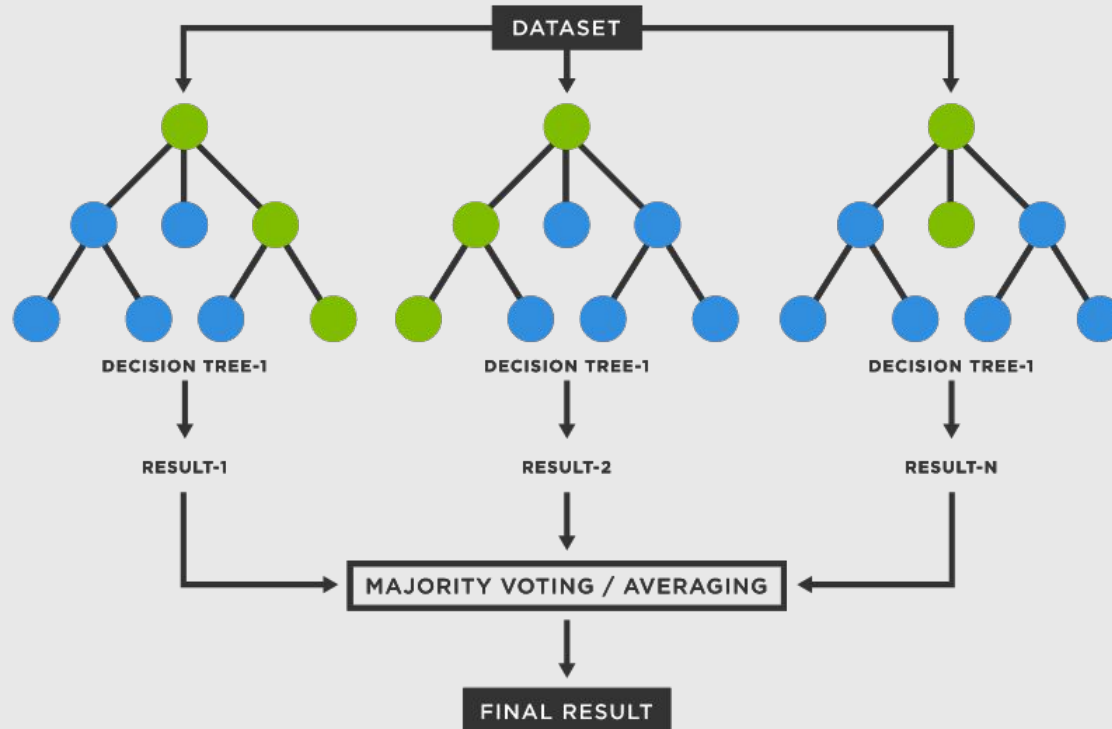
# DECISION TREE

Train Accuracy	Test Accuracy	Precision	Recall	F1 Score
91.92%	91.95%	0	0	0

	Predict Safe (0)	Predict Fraud (1)
Actual Safe (0)	56553	<b>0</b>
Actual Fraud (1)	<b>4950</b>	0

True Positive Rate	0%
False Positive Rate	0%

# RANDOM FOREST




# RANDOM FOREST

Train Accuracy	Test Accuracy	Precision	Recall	F1 Score
98.40%	91.87%	0.2298	0.0040	0.0079

	Predict Safe (0)	Predict Fraud (1)
Actual Safe (0)	56486	<b>67</b>
Actual Fraud (1)	<b>4930</b>	20

True Positive Rate	0.40%
False Positive Rate	0.11%

# K-NEAREST NEIGHBORS




# K-NEAREST NEIGHBORS

Train Accuracy	Test Accuracy	Precision	Recall	F1 Score
92.70%	90.29%	0.1308	0.0365	0.0571

	Predict Safe (0)	Predict Fraud (1)
Actual Safe (0)	55351	<b>1202</b>
Actual Fraud (1)	<b>4769</b>	181

True Positive Rate	3.65%
False Positive Rate	2.12%

# NAIVE BAYES



# NAIVE BAYES

Train Accuracy	Test Accuracy	Precision	Recall	F1 Score
91.30%	91.40%	0.0687	0.0054	0.0101

	Predict Safe (0)	Predict Fraud (1)
Actual Safe (0)	56187	<b>366</b>
Actual Fraud (1)	<b>4923</b>	27

True Positive Rate	0.54%
False Positive Rate	0.64%




# MODEL COMPARISON


	model	train_acc	test_acc	precision	recall	f1_score
1	Decision Tree	0.919210	0.919516	0.000000	0.000000	0.000000
2	Random Forest	0.984074	0.918752	0.229885	0.004040	0.007941
3	K-Nearest Neighbours	0.927096	0.902915	0.130875	0.036566	0.057161
4	Gaussian Naive Bayes	0.913023	0.914004	0.068702	0.005455	0.010107

# MODEL COMPARISON


- Highest test accuracy: **Decision Tree Classifier**
  - **91.95%**
  - Predicts all customers as safe (target=0)
  - Useless at detecting if customer will default
- Highest Recall (TPR): **K-Nearest Neighbors**
  - **3.65%**
  - Extremely low detection rate
  - Takes extremely long time to predict




# IMBALANCED TRAINING DATASET




# SYNTHETIC MINORITY OVER-SAMPLING TECHNIQUE (SMOTE)



Original Dataset




Generating Samples



Resampled Dataset

# RESAMPLED TRAINING DATASET



# CORRELATION (SMOTE)



**TARGET**

**OWNING A FLAT  
OWNING A CAR**

```
FLAG_OWN_CAR          -0.278647
FLAG_OWN_REALTY       -0.210410
DAYS_EMPLOYED         0.142916
NAME_CONTRACT_TYPE    -0.127142
DAYS_BIRTH            0.124180
AMT_GOODS_PRICE       -0.065508
AMT_INCOME_TOTAL      -0.051444
AMT_CREDIT            -0.044920
NAME_INCOME_TYPE      0.027394
OCCUPATION_TYPE       -0.023302
OBS_60_CNT_SOCIAL_CIRCLE -0.019977
OBS_30_CNT_SOCIAL_CIRCLE -0.018194
NAME_TYPE_SUITE       -0.016024
CNT_CHILDREN          0.008865
AMT_ANNUITY           0.003912
SK_ID_CURR            0.001456
DEF_30_CNT_SOCIAL_CIRCLE NaN
DEF_60_CNT_SOCIAL_CIRCLE NaN
Name: TARGET, dtype: float64
The most correlated factor with TARGET is 'FLAG_OWN_CAR', with a correlation value of -0.2786.
```

# Model Comparison - SMOTE

Rank	model	train_acc	test_acc	precision	recall	f1_score
4	Decision Tree	0.725909	0.692487	0.089975	0.309495	0.139418
1	Random Forest	0.993763	0.897013	0.128755	0.048485	0.070443
3	K-Nearest Neighbours	0.926815	0.724013	0.097429	0.293939	0.146349
2	Gaussian Naive Bayes	0.497342	0.893339	0.063922	0.023838	0.034726

# Model Comparison - SMOTE

- Sacrifice accuracy to improve model robustness
  - Decreased Accuracy
  - Increased Recall / F1 Score
- Highest test accuracy: **Random Forest Classifier**
  - **89.70% Accuracy**
  - Low recall: **4.84%**
- Highest Recall (TPR): **Decision Tree Classifier**
  - **30.94% Recall**
  - Low accuracy: **69.24%**



# Model Comparison - SMOTE

	model	train_acc	test_acc	precision	recall	f1_score
3	K-Nearest Neighbours	0.926815	0.724013	0.097429	0.293939	0.146349

BEST OVERALL MODEL: **K-Nearest Neighbours**

**GOOD**  
COMPROMISE **BETWEEN**

ACCURACY

&

RECALL





# CONCLUSION

# INSIGHTS & TAKEAWAYS

## 1 EDA

Understand overview of data



## 3 SMOTE

**REDUCES** ACCURACY



**BUT**

**INCREASES** PERFORMANCE



## 2 FEATURE SELECTION

Reduce training time



**AND**

Model performance not affected



# RECOMMENDATIONS

## INSIGHTS & RECOMMENDATIONS

Contributing Factors



DAYS\_EMPLOYED & DAYS\_BIRTH

## OUTCOME

SMOTE



Address class **imbalance** in dataset  
– under represented –



**INTRODUCE**

**NOISE  
OVERFITTING**

THANK YOU

