

CREDIT FRAUD DETECTION

SC1015 PRESENTATION

By: Jerald (U2221612D)
Bryan (U2222849J)
Jonathan (U2223191G)

TABLE OF CONTENTS

01

DATA SET USED

Motivations

02

DATA RELATED

Visualisation and Cleaning

03

MACHINE LEARNING

Decision Tree, Random Forest,
K-Nearest Neighbours, Naive Bayes

04

OUTCOME

Our Thoughts



DATA SET

MOTIVATION

- Cashless payments are **increasingly common** with the push of such technology
- To **minimise** the **risk** of lending money to customers who may not pay back their loans.
- To understand which **factors** affect people from defaulting credit card loans.



PAYNOW




73%

of Singaporeans own at least 1 credit card in 2023

DATA SET



application_data
columns_description
previous_application



FEATURE SELECTION

19/124 Columns in application_data

TARGET | NAME_CONTRACT_TYPE |
FLAG_OWN_CAR | FLAG_OWN_REALTY |
AMT_CREDIT | AMT_GOODS_PRICE, etc.


PROBLEM DEFINITION

"To find the **highest correlation** of credit card defaulters and to **predict** loan default risk of defaulters."



DATA CLEANING

SPLITTING THE DATA



NUMERICAL

- Finding **NULL** values within dataset
- Fill in missing values with **MEDIAN**

```
[ ] refined_app.isnull().sum()
```

```
SK_ID_CURR      0
TARGET          0
NAME_CONTRACT_TYPE 0
FLAG_OWN_CAR    0
FLAG_OWN_REALTY 0
CNT_CHILDREN    0
AMT_INCOME_TOTAL 0
AMT_CREDIT      0
```

```
AMT_ANNUITY      12
AMT_GOODS_PRICE  278
NAME_TYPE_SUITE  1292
NAME_INCOME_TYPE 0
DAYS_BIRTH       0
DAYS_EMPLOYED    0
OCCUPATION_TYPE  96391
OBS_30_CNT_SOCIAL_CIRCLE 1021
DEF_30_CNT_SOCIAL_CIRCLE 1021
OBS_60_CNT_SOCIAL_CIRCLE 1021
DEF_60_CNT_SOCIAL_CIRCLE 1021
dtype: int64
```

CATEGORICAL



- Finding **empty values** within dataset
- Fill in missing values with **NULL**


```
.isnull().sum()
```

FILLING MISSING VALUES

NUMERICAL

6/13 columns with missing data

```
def median_impute(df,col):  
    return df[col].fillna(df[col].median())  
  
refined_num['AMT_ANNUITY'] = median_impute(refined_num,'AMT_ANNUITY')
```



Filling in with **MEDIAN**

CATEGORICAL

2/6 columns with missing data

NAME_TYPE_SUITE: Fill with “*Unaccompanied*”

OCCUPATION_TYPE: Fill with “*NILL*”

.fillna

REMOVING DUPLICATES

- Ensure each record represents a unique observation
- Affects data analysis:
 - Skewness of results of statistical analysis
 - Increasing size of data set unnecessarily
 - Overfitting in machine learning



```
[ ] df = refined_app.copy()
    df.drop_duplicates(inplace=True)
    print('The amount of frauds in df before dropping duplicates:', len(refined_app[refined_app['TARGET'] == 1]))
    print('The amount of frauds in df after dropping duplicates:', len(df[df['TARGET'] == 1]))
```

```
The amount of frauds in df before dropping duplicates: 24825
The amount of frauds in df after dropping duplicates: 24825
```



DATA

VISUALISATION

NUMERICAL DATA

Target (whether individuals defaulted)




Heavily skewed - more 0s than 1s

Repercussions:

1. Biased model
2. Poor performance of minority class
3. Misleading evaluation metrics

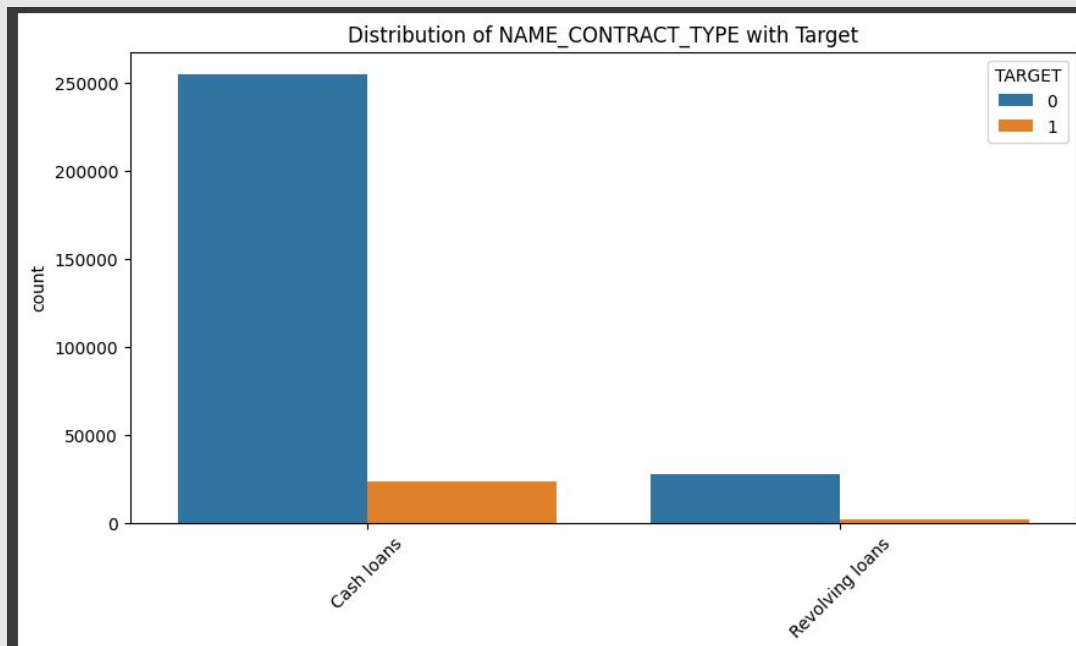
Possible solutions:

1. Undersampling of majority class
2. Oversampling of minority class (SMOTE)
3. Cost-sensitive learning
4. Ensemble method



CATEGORICAL DATA

UNDERSAMPLING



CORRELATION (RAW DATA)



TARGET

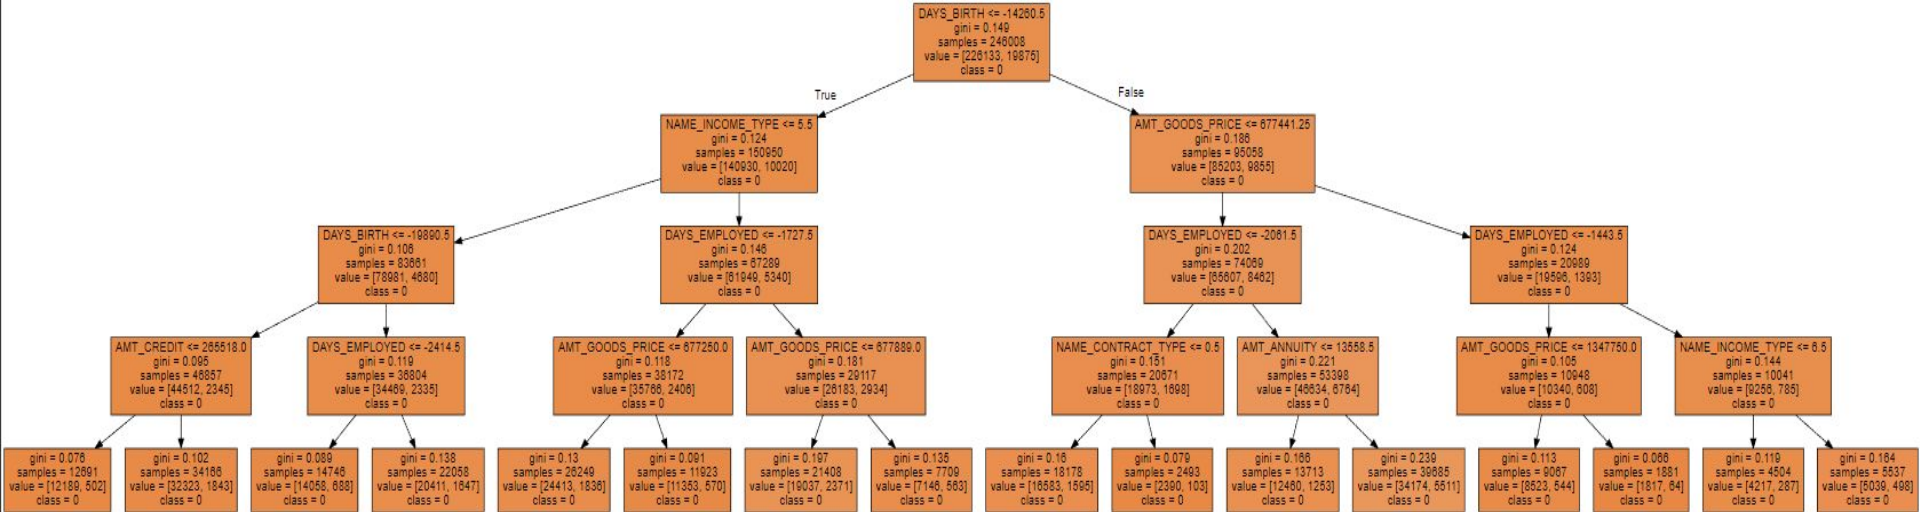
DAYS EMPLOYED
DAYS BIRTH

```
DAYS_EMPLOYED      0.071695
DAYS_BIRTH          0.063464
AMT_GOODS_PRICE    -0.029617
AMT_INCOME_TOTAL   -0.024722
AMT_CREDIT          -0.020198
AMT_ANNUITY         0.015305
CNT_CHILDREN        0.010850
OBS_60_CNT_SOCIAL_CIRCLE -0.009135
SK_ID_CURR          0.007665
OBS_30_CNT_SOCIAL_CIRCLE -0.007648
DEF_30_CNT_SOCIAL_CIRCLE      NaN
DEF_60_CNT_SOCIAL_CIRCLE      NaN
Name: TARGET, dtype: float64
The most correlated factor with TARGET is 'DAYS_EMPLOYED', with a correlation value of 0.0717.
Number of frauds (TARGET = 1) after dropping outliers: 1597
```

A black icon on the left side of the text, depicting a stylized brain with circuit-like lines extending from its right side, symbolizing machine learning or artificial intelligence.

MACHINE LEARNING

DECISION TREE



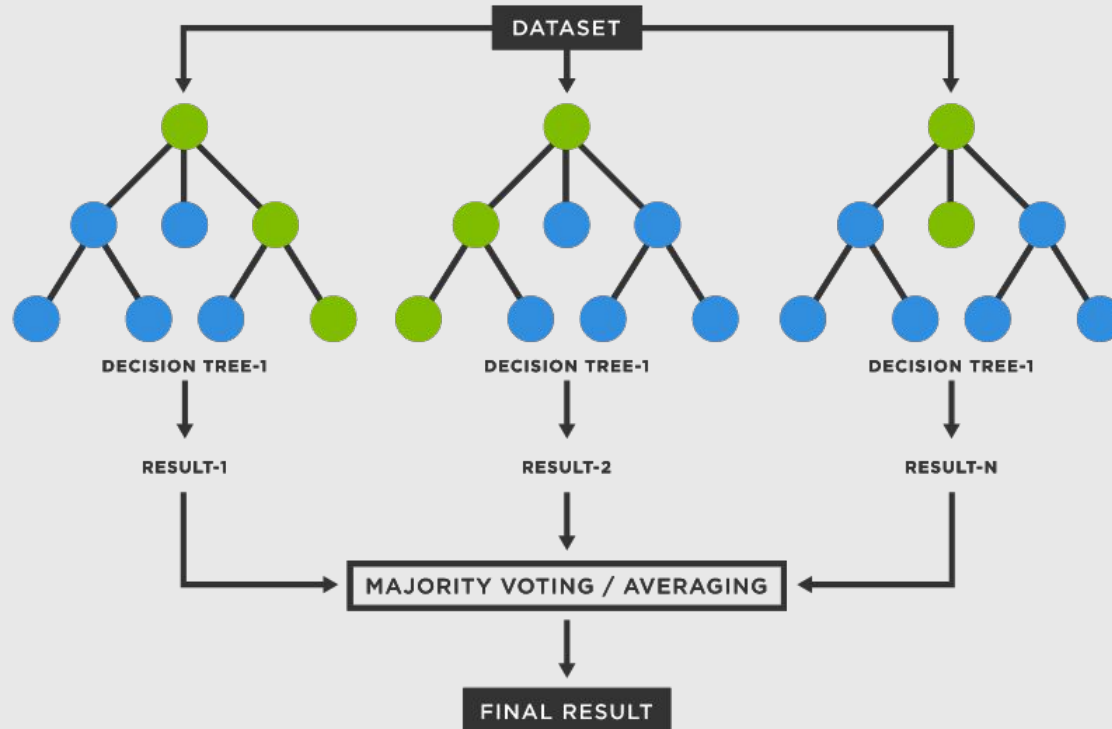
DECISION TREE

Train Accuracy	Test Accuracy	Precision	Recall	F1 Score
91.92%	91.95%	0	0	0

	Predict Safe (0)	Predict Fraud (1)
Actual Safe (0)	56553	0
Actual Fraud (1)	4950	0

True Positive Rate	0%
False Positive Rate	0%

RANDOM FOREST




RANDOM FOREST

Train Accuracy	Test Accuracy	Precision	Recall	F1 Score
98.40%	91.87%	0.2298	0.0040	0.0079

	Predict Safe (0)	Predict Fraud (1)
Actual Safe (0)	56486	67
Actual Fraud (1)	4930	20

True Positive Rate	0.40%
False Positive Rate	0.11%

K-NEAREST NEIGHBORS




K-NEAREST NEIGHBORS

Train Accuracy	Test Accuracy	Precision	Recall	F1 Score
92.70%	90.29%	0.1308	0.0365	0.0571

	Predict Safe (0)	Predict Fraud (1)
Actual Safe (0)	55351	1202
Actual Fraud (1)	4769	181

True Positive Rate	3.65%
False Positive Rate	2.12%

NAIVE BAYES



NAIVE BAYES

Train Accuracy	Test Accuracy	Precision	Recall	F1 Score
91.30%	91.40%	0.0687	0.0054	0.0101

	Predict Safe (0)	Predict Fraud (1)
Actual Safe (0)	56187	366
Actual Fraud (1)	4923	27


True Positive Rate	0.54%
False Positive Rate	0.64%

MODEL COMPARISON


	model	train_acc	test_acc	precision	recall	f1_score
1	Decision Tree	0.919210	0.919516	0.000000	0.000000	0.000000
2	Random Forest	0.984074	0.918752	0.229885	0.004040	0.007941
3	K-Nearest Neighbours	0.927096	0.902915	0.130875	0.036566	0.057161
4	Gaussian Naive Bayes	0.913023	0.914004	0.068702	0.005455	0.010107

MODEL COMPARISON


- Highest test accuracy: **Decision Tree Classifier**
 - **91.95%**
 - Predicts all customers as safe (target=0)
 - Useless at detecting if customer will default
- Highest Recall (TPR): **K-Nearest Neighbors**
 - **3.65%**
 - Extremely low detection rate
 - Takes extremely long time to predict




IMBALANCED TRAINING DATASET




SYNTHETIC MINORITY OVER-SAMPLING TECHNIQUE (SMOTE)



Original Dataset




Generating Samples



Resampled Dataset

RESAMPLED TRAINING DATASET



Model Comparison - SMOTE

Rank	model	train_acc	test_acc	precision	recall	f1_score
4	Decision Tree	0.725909	0.692487	0.089975	0.309495	0.139418
1	Random Forest	0.993763	0.897013	0.128755	0.048485	0.070443
3	K-Nearest Neighbours	0.926815	0.724013	0.097429	0.293939	0.146349
2	Gaussian Naive Bayes	0.497342	0.893339	0.063922	0.023838	0.034726

Model Comparison - SMOTE

- Sacrifice accuracy to improve model robustness
 - Decreased Accuracy
 - Increased Recall / F1 Score
- Highest test accuracy: **Random Forest Classifier**
 - **89.70% Accuracy**
 - Low recall: **4.84%**
- Highest Recall (TPR): **Decision Tree Classifier**
 - **30.94% Recall**
 - Low accuracy: **69.24%**

Model Comparison - SMOTE

	model	train_acc	test_acc	precision	recall	f1_score
3	K-Nearest Neighbours	0.926815	0.724013	0.097429	0.293939	0.146349

BEST OVERALL MODEL: **K-Nearest Neighbours**

GOOD
COMPROMISE **BETWEEN**

ACCURACY

&

RECALL





CONCLUSION

INSIGHTS & TAKEAWAYS

1 EDA

Understand overview of data



3 SMOTE

REDUCES ACCURACY



BUT

INCREASES PERFORMANCE



2 FEATURE SELECTION

Reduce training time



AND

Model performance not affected



RECOMMENDATIONS

INSIGHTS & RECOMMENDATIONS

Contributing Factors



DAYS_EMPLOYED & DAYS_BIRTH

OUTCOME

SMOTE



Address class **imbalance** in dataset
– under represented –



INTRODUCE

**NOISE
OVERFITTING**

THANK YOU

